

Allen School of Computer Science & Engineering
University of Washington
3800 E Stevens Way NE, CSE2 233
Seattle, WA, USA 98195

(760) 689-2466 
[@cquanze](https://twitter.com/cquanze) 
cqz@cs.washington.edu 
homes.cs.washington.edu/~cqz 

QUAN ZE CHEN

APPOINTMENTS

2023–2025 **University of Washington**, Seattle, WA
Postdoctoral Scholar, Computer Science & Engineering
Advisors: Amy X. Zhang, Franziska Roesner

EDUCATION

2018–2023 **University of Washington**, Seattle, WA
Ph.D. Computer Science & Engineering
Advisor: Amy X. Zhang

2016–2018 **University of Washington**, Seattle, WA
M.S. Computer Science & Engineering

2012–2016 **University of Pennsylvania**, Philadelphia, PA
B.S.E. Networked and Social Systems Engineering, Summa Cum Laude

PUBLICATIONS

Conference and Journal Papers

- [c.17] **Quan Ze Chen**, Amy X. Zhang. Case Law Grounding: Using Precedents to Align Decision-Making for Humans and AI. *ACM Conference On Collective Intelligence (CI '25)*. La Jolla, CA. August 2025.
- [c.16] **Quan Ze Chen**, K.J. Kevin Feng, Chan Young Park, Amy X. Zhang. SPICA: Retrieving Scenarios for Pluralistic In-Context Alignment. *ACL Findings 2025*. Vienna, Austria. July 2025.
- [c.15] Aviv Ovadya, Kyle Redman, Oliver Smith, Luke Thorburn, **Quan Ze Chen**, Flynn Devine, Andrew Konya, Smitha Milli, Manon Revel, K. J. Kevin Feng, Amy X. Zhang, Bilva Chandra, Michiel A. Bakker, Atoosa Kasirzadeh. Position: Democratic AI is Possible. The Democracy Levels Framework Shows How It Could Work. *The 42nd International Conference on Machine Learning*. Vancouver, Canada. July 2025.
- [c.14] Tzu-Sheng Kuo, **Quan Ze Chen**, Amy X. Zhang, Jane Hsieh, Haiyi Zhu, Kenneth Holstein. PolicyCraft: Supporting Collaborative and Participatory Policy Design through Case-Grounded Deliberation. *ACM Conference on Human Factors in Computing Systems (CHI '25)*. Yokohama, Japan. April 2025.
- [c.13] Leijie Wang, Kathryn Yurechko, Pranati Dani, **Quan Ze Chen**, Amy X. Zhang. End User Authoring of Personalized Content Classifiers: Comparing Example Labeling, Rule Writing, and LLM Prompting. *ACM Conference on Human Factors in Computing Systems (CHI '25)*. Yokohama, Japan. April 2025.
- [c.12] Inyoung Cheong, King Xia, K.J. Kevin Feng, **Quan Ze Chen**, Amy X. Zhang. (A)I Am Not a Lawyer, But...: Engaging Legal Experts towards Responsible LLM Policies for Legal Advice. *ACM Conference on Fairness, Accountability, and Transparency (FAccT '24)*. Rio de Janeiro, Brazil. June 2024.

- [c.11] Kianna Bolante, Kevin Chen, **Quan Ze Chen**, Amy X. Zhang. Bringing Social Computing to Secondary School Classrooms. *ACM Special Interest Group on Computer Science Education Technical Symposium (SIGCSE TS 2024)*. Portland, OR. March 2024.
- [c.10] Andre Ye, **Quan Ze Chen**, Amy X. Zhang. Confidence Contours: Uncertainty-Aware Annotation for Medical Semantic Segmentation. *AAAI Conference on Human Computation and Crowdsourcing (HCOMP '23)*. Delft, Netherlands. November 2023. **Honorable Mention**
- [c.9] **Quan Ze Chen**, Amy X. Zhang. Judgment Sieve: Reducing Uncertainty in Group Judgments through Interventions Targeting Ambiguity versus Disagreement. *ACM Conference On Computer-Supported Cooperative Work And Social Computing (CSCW '23)*. Minneapolis, MN. October 2023.
- [c.8] Shagun Jhaver, Alice Zhang, **Quan Ze Chen**, Nikhila Natarajan, Ruotong Wang, Amy X. Zhang. Personalizing Content Moderation on Social Media: User Perspectives on Moderation Choices, Interface Design, and Labor. *ACM Conference On Computer-Supported Cooperative Work And Social Computing (CSCW '23)*. Minneapolis, MN. October 2023.
- [c.7] Teanna Barrett, **Quan Ze Chen**, Amy X. Zhang. Skin Deep: Investigating Subjectivity in Skin Tone Annotations for Computer Vision Benchmark Datasets. *ACM Conference on Fairness Accountability and Transparency (FAccT '23)*. Chicago, IL. June 2023.
- [c.6] Shagun Jhaver, **Quan Ze Chen**, Detlef Knauss, Amy X. Zhang. Designing Word Filter Tools for Creator-led Comment Moderation. *ACM Conference on Human Factors in Computing Systems (CHI '22)*. New Orleans, LA. April 2022.
- [c.5] **Quan Ze Chen**, Tobias Schnabel, Besmira Nushi, Saleema Amershi. HINT: Integration Testing for AI-based features with Humans in the Loop. *ACM Conference on Intelligent User Interfaces (IUI '22)*. Virtual. March 2022.
- [c.4] **Quan Ze Chen**, Daniel S. Weld, Amy X. Zhang. Goldilocks: Consistent Crowdsourced Scalar Annotations with Relative Uncertainty. *Proceedings of the ACM on Human-Computer Interaction (CSCW '21)*. Virtual. October 2021.
- [c.3] **Quan Ze Chen**, Jonathan Bragg, Lydia B. Chilton, Daniel S. Weld. Cicero: Multi-Turn, Contextual Argumentation for Accurate Crowdsourcing. *ACM Conference on Human Factors in Computing Systems (CHI '19)*. Glasgow, United Kingdom. May 2019.
- [c.1] Wei Xu, Courtney Napoles, Ellie Pavlick, **Quan Ze Chen**, Chris Callison-Burch. Optimizing Statistical Machine Translation for Simplification. *Transactions of the Association for Computational Linguistics (ACL '16)*. Berlin, Germany. August 2016.

Posters and Work-in-Progress Papers

- [p.6] Aviv Ovadya, Kyle Redman, Oliver Smith, Luke Thorburn, **Quan Ze Chen**, Flynn Devine, Andrew Konya, Smitha Milli, Manon Revel, K. J. Kevin Feng, Amy X. Zhang, Bilva Chandra, Michiel A. Bakker, Atoosa Kasirzadeh. The Democracy Levels Framework: Towards Integrating Democratic Infrastructure into the AI Ecosystem. *First Workshop on Sociotechnical AI Governance: Opportunities and Challenges for HCI (STAIG Workshop at CHI '25)*. Yokohama, Japan. Apr 2025.
- [p.5] K. J. Kevin Feng, **Quan Ze Chen**, Inyoung Cheong, Amy X. Zhang. Policy Prototyping for LLMs: Pluralistic Alignment via Interactive and Collaborative Policymaking. *Workshop on Bidirectional Human-AI Alignment (Bi-Align Workshop at ICLR '25)*. Singapore. Apr 2025.
- [p.4] Andrew Konya, Aviv Ovadya, Kevin Feng, **Quan Ze Chen**, Lisa Schirch, Colin Irwin, Amy X Zhang. Chain of Alignment: Integrating Public Will with Expert Intelligence for Language

Model Alignment. *Pluralistic Alignment, Conference on Neural Information Processing Systems (Pluralistic Alignment Workshop at NeurIPS '23)*. Vancouver, Canada. Dec 2024.

- [p.3] K. J. Kevin Feng, **Quan Ze Chen**, Inyoung Cheong, King Xia, Amy X. Zhang. Case Repositories: Towards Case-Based Reasoning for AI Alignment. *First Workshop on AI meets Moral Philosophy and Moral Psychology, Conference on Neural Information Processing Systems (MP2 Workshop at NeurIPS '23)*. New Orleans, LA. Dec 2023.
- [p.2] John Akers, Gagan Bansal, Gabriel Cadamuro, Christine Chen, **Quan Ze Chen**, Lucy Lin, Phoebe Mulcaire, Rajalakshmi Nandakumar, Matthew Rockett, Lucy Simko, John Toman, Tongshuang Wu, Eric Zeng, Bill Zorn, Franziska Roesner. Technology-enabled disinformation: Summary, lessons, and recommendations. *arXiv 2018*. Digital Preprint.
- [p.1] **Quan Ze Chen**, Chenyang Lei, Wei Xu, Ellie Pavlick, Chris Callison-Burch. Poetry of the Crowd: A Human Computation Algorithm to Convert Prose into Rhyming Verse *AAAI Conference on Human Computation and Crowdsourcing (HCOMP '14)*. Pittsburgh, PA. Sept 2014.

Doctoral Symposium

- [ds.1] **Quan Ze Chen**. Disentangling Inherent Ambiguity and Disagreement in Crowdsourced Annotation. *AAAI Conference on Human Computation and Crowdsourcing (HCOMP '21)*. Virtual. Nov 2021.

INDUSTRY EXPERIENCE AND INTERNSHIPS

- 2025-2026 Researcher, **AI & Democracy Foundation**, Covina, CA (Remote)
Mentor: Aviv Ovadya
- Summer 2021 Applied Scientist Intern, **SageMaker Ground Truth, Amazon AI**, Palo Alto, CA
Mentors: Mike Schaeckermann, Matt Lease, Erran Li
- Spring 2020 Research Intern, **Human-AI eXperiences (HAX) Team, Microsoft Research**, Redmond, WA
Mentors: Besmira Nushi, Tobias Schnabel, Saleema Amershi
- Spring 2019 Research Intern, **Semantic Scholar, AI2**, Seattle, WA
Mentor: Dan Weld
- Summer 2015 Software Engineer Intern, **Google Inc.**, Los Angeles, CA
- 2010-2012 High School Research Intern, **Institute of Microbiology, Chinese Academy of Sciences**, Beijing, China
Mentor: Baoli Zhu

GRANTS

- 2023 *Democratic Inputs to AI, OpenAI Inc.*, Case Law for AI Policy: Democratically Aligning AI with Human Preferences through Case Repositories

TEACHING

Courses

- Autumn 2022 Teaching Assistant, **CSE 440: Introduction To HCI**, UW CSE
Graduate TA for undergraduate special topics course.
- Spring 2021 Teaching Assistant, **CSEP 590B: Enterprise Chatbots**, UW CSE

	Graduate TA for MS level special topics course.
Winter 2021	Teaching Assistant, CSE 415: Introduction to Artificial Intelligence , UW CSE Graduate TA for non-majors AI intro course.
Winter 2019	Teaching Assistant, CSE 473: Introduction to Artificial Intelligence , UW CSE Graduate TA for majors AI intro course.
Spring 2019	Teaching Assistant, CSEP 573: Introduction to Artificial Intelligence , UW CSE Graduate TA for MS AI intro course.
2013 - 2016	Teaching Assistant, CIS 121: Data Structures and Algorithms , University of Pennsylvania Undergraduate TA, Head TA for undergraduate intro CS course.

INVITED TALKS

Nov 2023	Colloquium , University of Texas, Austin, TX Working with Uncertainty to Align Humans and AI on Socially-Constructed Concepts
Sept 2023	Democratic Inputs to AI , OpenAI Inc., San Francisco, CA Case Law for AI Policy
July 2023	ASE Events Invited Talk , Apple Inc., Seattle, WA Tools and Workflows for Aligning Machine Learning Systems with Human Uncertainty
Aug 2022	Stanford HCI Lunch , Stanford University, Virtual Understanding and Addressing Uncertainty of the Crowd
June 2022	Expertise@Scale Salon , Carnegie Mellon University, Virtual Understanding and Addressing Uncertainty of the Crowd

SERVICE

Reviewer

Conferences: CHI '20-'24, CSCW '21-'25, WWW '20, CI '25

Journals: JAIR

Workshops: CHI '22-'23 TRAIT, CHI '24 TREW, NeurIPS '24 Pluralistic Alignment, CHI '25 STAIG

STUDENT ADVISING

Undergraduate Students

2023	Kathryn Yurechko
2022	Andre Ye (Now PhD Student at MIT) Teanna Barrett (Now PhD Student at UW) Kianna Bolante